

Importing required libraries

these are useful as they are needed to run all the functions in our notebook

```
In [1]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.preprocessing import image
from tensorflow.keras.optimizers import RMSprop
from tensorflow.keras.layers import Conv2D, Flatten, Dense, Dropout, Input
from tensorflow.keras.models import Model
import tensorflow as tf
import matplotlib.pyplot as plt
from skimage import io
from tqdm import tqdm
import numpy as np
import os
import shutil
import random
import pandas as pd
import cv2 as cv
# from PIL import Image
```

Generating xtrain and ytrain

Here we are generating xtrain, for that we have to go through the following steps -

- 1) locate to the train folder
- 2) convert the images to grayscale, we are converting to grayscale as other colors aren't present in our dataset, black and white are the colors available
- 3) resize the image to 28,28. We are doing this as the standard MNIST dataset is also in 28,28. also cause 28x28 pixels are enough to train our model
- 4) appending the image pixels value and label in an array to finally get our xtrain

Printing a counter in the loop so we know how much images have been processed yet

```
In [2]: train_dir = os.listdir(r'E:\Internship\IIITD MIDAS\Dataset\Dataset\train')
root_dir = r'E:\Internship\IIITD MIDAS\Dataset\Dataset'
X = []
f=0
for cls in train_dir:
    src = root_dir + '\\train\\' + cls
    allFileNames = os.listdir(src)
    a = cls
    a = a[-2:]
    p = int(a)
    p
    # print(cls)
    # print(src + '/' + allFileNames)

    for img in allFileNames:
        # ytrain.append(p)
        IMG_LOC=src+'\\'+img
        # image = io.imread(IMG_LOC)
        image=cv.imread(IMG_LOC, cv.IMREAD_GRAYSCALE)
        image=cv.resize(image,(28,28))
        print(f)
        f=f+1
        # testing_data.append([np.array(img),img_num])
        X.append([np.array(image),np.array(p)])
    # print(X)
    # print(ytrain)
    # print(image)
```

```
2088
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2100
2101
2102
2103
2104
2105
2106
2107
```

Reshaping the array X and dividing it in Xtrain and Ytrain. We need xtrain and y train to train our model

```
In [3]: x_train = np.array([i[0] for i in X]).reshape(-1,28,28,1)
y_train = np.array([i[1] for i in X])
```

Printing the shape of Xtrain and Ytrain, to verify if its same as what we need in our model

```
In [4]: print(x_train.shape)
print(y_train.shape)
```

```
(2108, 28, 28, 1)
(2108,)
```

Generating Xtest and Ytest

Following the same steps as we did for xtrain and ytrain

```
In [5]: train_dir = os.listdir(r'E:\Internship\IIITD MIDAS\Dataset\Dataset\test')
# img=cv.imread("./content/drive/MyDrive/IIITD/data/Sample027/img027-004.png", cv.IMREAD_GRAYSCALE)
root_dir = r'E:\Internship\IIITD MIDAS\Dataset\dataset'
# os.makedirs(root_dir + '/train')

y = []
# ytrain = []
f=0
for cls in train_dir:
    src = root_dir+ '\\\\test\\\\' + cls
    allFileNames = os.listdir(src)
    a = cls
    a = a[-2:]
    p = int(a)
    p
    # print(cls)
    # print(src + '/' + allFileNames)
    for img in allFileNames:
        # ytrain.append(p)
        IMG_LOC=src+'\\\\'+img
        # image = io.imread(IMG_LOC)
        image=cv.imread(IMG_LOC, cv.IMREAD_GRAYSCALE)
        image=cv.resize(image,(28,28))
        # testing_data.append([np.array(img),img_num])
        y.append([np.array(image),np.array(p)])
        print(f)
        f=f+1
    # print(y)
    # print(ytrain)
    # print(image)

353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
```

```
In [6]: x_test = np.array([i[0] for i in y]).reshape(-1,28,28,1)
y_test = np.array([i[1] for i in y])
```

```
In [7]: print(x_test.shape)
print(y_test.shape)
print(x_train.shape)
print(y_train.shape)
```

```
(372, 28, 28, 1)
(372,)
(2108, 28, 28, 1)
(2108,)
```

Data Normalization

As the CNN algorithm converge faster on the [0..1] data than on [0..255]

```
In [8]: x_train, x_test=x_train/255.0, x_test/255.0
y_train, y_test=y_train.flatten(), y_test.flatten()
```

```
In [9]: y_train
```

```
Out[9]: array([ 1,  1,  1, ..., 62, 62, 62])
```

Label encoding

We need to give the y test and train in one hot encoded form

used this article for Label encoding - <https://www.geeksforgeeks.org/ml-label-encoding-of-datasets-in-python/>

it basically converts all the labels from 0 to the end. so that it becomes easier in one hot encoding.

```
In [10]: # Import Label encoder
from sklearn import preprocessing

# label_encoder object knows how to understand word labels.
label_encoder = preprocessing.LabelEncoder()

# Encode Labels in column 'species'.
y_train=label_encoder.fit_transform(y_train)
y_train

Out[10]: array([ 0,  0,  0, ..., 61, 61, 61], dtype=int64)
```

Used the `to_categorical` function to convert the data to one hot encoded form

```
In [11]: from numpy import array
from numpy import argmax
from keras.utils import to_categorical
# define example
data = array(y_train)
print(data)
# one hot encode
encoded = to_categorical(data)
print(encoded)
y_train = encoded

[[ 0.  0.  0. ... 61. 61. 61.]
 [1.  0.  0. ... 0.  0.  0.]
 [1.  0.  0. ... 0.  0.  0.]
```

```
[1. 0. 0. ... 0. 0. 0.]  
...  
[0. 0. 0. ... 0. 0. 1.]  
[0. 0. 0. ... 0. 0. 1.]
```

Same for Y test

```
In [12]: label_encoder = preprocessing.LabelEncoder()  
y_test = label_encoder.fit_transform(y_test)
```

```
In [13]: y_test
```

```
Out[13]: array([ 0,  0,  0,  0,  0,  1,  1,  1,  1,  1,  1,  1,  2,  2,  2,  2,  2,  
    2,  3,  3,  3,  3,  3,  4,  4,  4,  4,  4,  4,  4,  5,  5,  5,  5,  
    5,  5,  6,  6,  6,  6,  6,  7,  7,  7,  7,  7,  7,  8,  8,  8,  
    8,  8,  9,  9,  9,  9,  9,  9,  9,  9,  9,  9,  9,  10, 10, 10, 10,  
    10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 11, 11,  
    11, 11, 11, 11, 12, 12, 12, 12, 12, 12, 12, 13, 13, 13, 13, 13, 13, 13, 14,  
    14, 14, 14, 14, 14, 15, 15, 15, 15, 15, 15, 15, 16, 16, 16, 16, 16, 16,  
    17, 17, 17, 17, 17, 17, 18, 18, 18, 18, 18, 18, 18, 19, 19, 19, 19, 19,  
    19, 20, 20, 20, 20, 20, 21, 21, 21, 21, 21, 21, 21, 22, 22, 22, 22,  
    22, 22, 23, 23, 23, 23, 23, 24, 24, 24, 24, 24, 24, 24, 25, 25, 25,  
    25, 25, 25, 26, 26, 26, 26, 26, 27, 27, 27, 27, 27, 27, 27, 28, 28,  
    28, 28, 28, 29, 29, 29, 29, 29, 29, 29, 29, 29, 29, 30, 30, 30, 30, 30,  
    31, 31, 31, 31, 31, 32, 32, 32, 32, 32, 32, 32, 33, 33, 33, 33, 33,  
    34, 34, 34, 34, 34, 35, 35, 35, 35, 35, 35, 35, 36, 36, 36, 36, 36,  
    36, 37, 37, 37, 37, 37, 38, 38, 38, 38, 38, 38, 38, 39, 39, 39, 39,  
    39, 39, 40, 40, 40, 40, 40, 40, 41, 41, 41, 41, 41, 41, 41, 42, 42, 42,  
    42, 42, 42, 43, 43, 43, 43, 43, 44, 44, 44, 44, 44, 44, 44, 45, 45,  
    45, 45, 45, 46, 46, 46, 46, 46, 46, 47, 47, 47, 47, 47, 47, 47, 48,  
    48, 48, 48, 48, 49, 49, 49, 49, 49, 49, 49, 49, 50, 50, 50, 50, 50, 50, 50,  
    51, 51, 51, 51, 51, 51, 52, 52, 52, 52, 52, 52, 52, 53, 53, 53, 53, 53,  
    53, 54, 54, 54, 54, 54, 55, 55, 55, 55, 55, 55, 55, 56, 56, 56, 56,  
    56, 56, 57, 57, 57, 57, 57, 58, 58, 58, 58, 58, 58, 58, 59, 59, 59,  
    59, 59, 59, 60, 60, 60, 60, 60, 60, 61, 61, 61, 61, 61, 61]
```

```
In [14]: data = array(y_test)  
print(data)  
# one hot encode  
encoded = to_categorical(data)  
print(encoded)  
y_test = encoded
```

```
[ 0  0  0  0  0  0  1  1  1  1  1  1  1  2  2  2  2  2  2  2  3  3  3  3  3  3  
 4  4  4  4  4  4  5  5  5  5  5  5  5  6  6  6  6  6  6  6  7  7  7  7  7  7  
 8  8  8  8  8  9  9  9  9  9  9  9  9  10 10 10 10 10 10 10 11 11 11 11 11  
12 12 12 12 12 12 13 13 13 13 13 13 13 14 14 14 14 14 14 14 15 15 15 15 15  
16 16 16 16 16 17 17 17 17 17 17 17 18 18 18 18 18 18 19 19 19 19 19 19  
20 20 20 20 20 21 21 21 21 22 22 22 22 22 23 23 23 23 23 23 23 23  
24 24 24 24 24 25 25 25 25 25 25 26 26 26 26 26 26 27 27 27 27 27 27  
28 28 28 28 28 29 29 29 29 29 29 30 30 30 30 30 30 31 31 31 31 31  
32 32 32 32 32 32 33 33 33 33 33 33 34 34 34 34 34 34 35 35 35 35 35  
36 36 36 36 36 36 37 37 37 37 37 37 38 38 38 38 38 38 39 39 39 39 39  
40 40 40 40 40 41 41 41 41 41 41 42 42 42 42 42 43 43 43 43 43 43  
44 44 44 44 44 45 45 45 45 45 45 46 46 46 46 46 47 47 47 47 47 47  
48 48 48 48 48 49 49 49 49 49 49 49 49 49 49 50 50 50 50 50 50 51 51 51 51  
52 52 52 52 52 52 53 53 53 53 53 53 54 54 54 54 54 54 55 55 55 55 55 55  
56 56 56 56 56 56 57 57 57 57 57 57 58 58 58 58 58 58 59 59 59 59 59 59  
60 60 60 60 60 60 61 61 61 61 61 61 61 61 61 61 61 61 61 61 61 61]
```

```
In [15]: print(x_test.shape)  
print(y_test.shape)  
print(x_train.shape)  
print(y_train.shape)
```

```
(372, 28, 28, 1)  
(372, 62)  
(2108, 28, 28, 1)  
(2108, 62)
```

SAVING CHECKPOINTS

We need to save the checkpoints of our model, in case any error occurs in between our training

I referred this article to save checkpoints - <https://towardsdatascience.com/checkpointing-deep-learning-models-in-keras-a652570b8de6>

```
In [16]: checkpoint_path = "E:\Internship\IIITD MIDAS\checkpoints\cp-{epoch:04d}.ckpt"  
cp_callback = tf.keras.callbacks.ModelCheckpoint(  
    filepath=checkpoint_path,  
    save_weights_only=True,  
    save_freq=200,  
    verbose=1)
```

Model

Importing required libraries

```
In [17]: from sklearn.metrics import confusion_matrix  
import tensorflow as tf  
  
from keras.utils import to_categorical # convert to one-hot-encoding  
from keras.models import Sequential  
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D  
from keras.optimizers import RMSprop  
from keras.preprocessing.image import ImageDataGenerator  
from keras.callbacks import ReduceLROnPlateau
```

Here we have created a function to create a CNN model by adding the following Layers

I have added 2 convolution layers, with filter size as 32. Padding as Same, and activation function as relu

one Max Pooling layer with pool size 2,2 and 1 Dropout to prevent overfitting

I have repeated the above architecture with a filter size of 64, as with more layers, the model might perform better

Finally I have added a flatten layer to flatten the input, so that I can feed it to the dense layer.

I have then added the dense layer with relu as activation function.

At last I have added another dense layer with softmax as activation function because my final output has 62 classes.

After this I have added the required optimizers and I have finally compiled the model

```
In [18]: def create_model():
    model = Sequential()

    model.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',
                    activation ='relu', input_shape = (28,28,1)))
    model.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',
                    activation ='relu'))
    model.add(MaxPool2D(pool_size=(2,2)))
    model.add(Dropout(0.25))

    model.add(Conv2D(filters = 64, kernel_size = (3,3),padding = 'Same',
                    activation ='relu'))
    model.add(Conv2D(filters = 64, kernel_size = (3,3),padding = 'Same',
                    activation ='relu'))
    model.add(MaxPool2D(pool_size=(2,2), strides=(2,2)))
    model.add(Dropout(0.25))

    model.add(Flatten())
    model.add(Dense(256, activation = "relu"))
    model.add(Dropout(0.5))
    model.add(Dense(62, activation = "softmax"))

    optimizer = RMSprop(lr=0.001,
                        rho=0.9,
                        epsilon=1e-08,
                        decay=0.0)

    model.compile(optimizer = optimizer ,
                  loss = "categorical_crossentropy",
                  metrics=["accuracy"])
    return model
```

Creating a Model

```
In [19]: model_ckpt=create_model()
```

Setting the number of epochs and batch size

```
In [20]: epochs = 300 # Turn epochs to 30 to get 0.9967 accuracy
batch_size = 86
```

Data Augmentation

In data augmentation, we add a few images to the training dataset, with a little variation such as rotation or whitening and flipping of images. We do so, so that our model does not overfit.

Getting over fit means our model will get so much used to our training dataset, that its prediction on new images will be very faulty

```
In [21]: datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range=10, # randomly rotate images in the range (degrees, 0 to 180)
    zoom_range = 0.1, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction of total height)
    horizontal_flip=False, # randomly flip images
    vertical_flip=False) # randomly flip images

datagen.fit(x_train)
```

We are finally running the model we have built, by giving all the attributes to the `model.fit_generator` function

```
In [22]: history = model_ckpt.fit_generator(datagen.flow(x_train,y_train, batch_size=batch_size),
                                         epochs = epochs,
                                         validation_data = (x_test,y_test),
                                         verbose = 2,
                                         steps_per_epoch=x_train.shape[0] // batch_size
                                         , callbacks=[cp_callback])

24/24 - 14s - loss: 0.3493 - accuracy: 0.8828 - val_loss: 0.5227 - val_accuracy: 0.8737
Epoch 293/300
24/24 - 13s - loss: 0.3144 - accuracy: 0.8881 - val_loss: 0.5409 - val_accuracy: 0.8629
Epoch 294/300
24/24 - 13s - loss: 0.3437 - accuracy: 0.8773 - val_loss: 0.5610 - val_accuracy: 0.8602
Epoch 295/300
24/24 - 15s - loss: 0.3729 - accuracy: 0.8744 - val_loss: 0.5968 - val_accuracy: 0.8710
Epoch 296/300
24/24 - 14s - loss: 0.3562 - accuracy: 0.8823 - val_loss: 0.5413 - val_accuracy: 0.8575
Epoch 297/300
24/24 - 14s - loss: 0.3401 - accuracy: 0.8828 - val_loss: 0.4674 - val_accuracy: 0.8710
Epoch 298/300
24/24 - 14s - loss: 0.3697 - accuracy: 0.8724 - val_loss: 0.4850 - val_accuracy: 0.8602
Epoch 299/300
24/24 - 14s - loss: 0.3971 - accuracy: 0.8615 - val_loss: 0.4789 - val_accuracy: 0.8602
Epoch 300/300

Epoch 00300: saving model to E:\Internship\IIITD MIDAS\checkpoints\cp-0300.ckpt
24/24 - 15s - loss: 0.3504 - accuracy: 0.8778 - val_loss: 0.4783 - val_accuracy: 0.8656
```

Calculating the final accuracy and loss

```
In [25]: loss,acc = model_ckpt.evaluate(x_test, y_test, verbose=2)
```

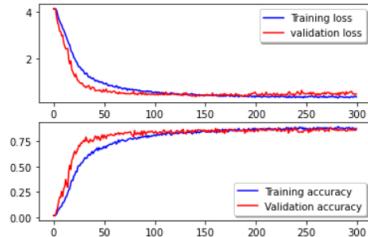
```
12/12 - 0s - loss: 0.4783 - accuracy: 0.8656
```

Evaluation

Plotting the loss and accuracy curves for training and validation

```
In [27]: fig, ax = plt.subplots(2,1)
ax[0].plot(history.history['loss'], color='b', label="Training loss")
ax[0].plot(history.history['val_loss'], color='r', label="validation loss", axes=ax[0])
legend = ax[0].legend(loc='best', shadow=True)

ax[1].plot(history.history['accuracy'], color='b', label="Training accuracy")
ax[1].plot(history.history['val_accuracy'], color='r', label="Validation accuracy")
legend = ax[1].legend(loc='best', shadow=True)
```



```
In [29]: model = model_ckpt
```

Plotting the Confusion matrix

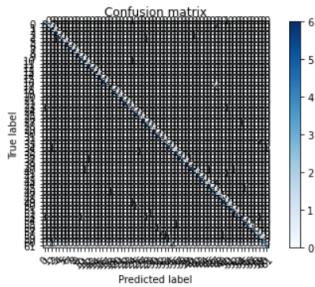
```
In [39]: def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

# Predict the values from the validation dataset
y_pred = model.predict(x_test)
# Convert predictions classes to one hot vectors
y_pred_classes = np.argmax(y_pred, axis = 1)
# Convert validation observations to one hot vectors
y_true = np.argmax(y_test, axis = 1)
# compute the confusion matrix
confusion_mtx = confusion_matrix(y_true, y_pred_classes)
# plot the confusion matrix
plot_confusion_matrix(confusion_mtx, classes = range(62))
```



here the confusion matrix is not clearly visible because of so many number of classes

```
In [ ]:
```

```
In [ ]:
```